

The University of Alabama in Huntsville
ECE Department
CPE 431 01, CPE 531 01/01R
Fall 2022
Parallel Processors from Client to Cloud

Due November 29, 2022, 1.0(10), 2.0(10), 3.0(10)

1.0 Consider the following piece of C code:

```
for (j = 2; j < 1000; j++)
    D[j] = D[j-1] + D[j-2];
```

The MIPS code corresponding to the above fragment is:

```

        addiu    $s2, $zero, 7992
        addiu    $s1, $zero, 16
Loop:   l.d      $f0, -16($s1)
        l.d      $f2, -8($s1)
        add.d    $f4, $f0, $f2
        s.d      $f4, 0($s1)
        addiu    $s1, $s1, 8
        bne     $s1, $s2, loop

```

Instructions have the following associated latencies (in cycles):

add.d	l.d	s.d	addiu
3	4	2	1

1.0.1 How many cycles does it take for all instructions in a single iteration of the above loop to execute?

1.0.2 When an instruction in a later iteration of a loop depends upon a data value produced in an earlier iteration of the same loop, we say that there is a loop-carried dependence between iterations of the loop. Identify the loop-carried dependences in the above code. Identify the dependent program variable and assembly-level registers. You can ignore the loop induction variable *j*.

2.0 Consider the following portions of two different programs running at the same time on four processors in a symmetric multi-core processor (SMP). Assume that before this code is run *w*, *x*, *y*, and *z* are -1, 4, 3 and 2 respectively and they are declared as int,

Core 1: *x* = *z* + *w*;

Core 2: *y* = *x* + *z*;

Core 3: *w* = *x***y* + 1;

Core 4: *z* = *x*/*w*;

What are all the possible resulting values of *w*, *x*, *y*, and *z*? For each possible outcome, explain how we might arrive at these values. You will need to examine all possible interleavings of instructions.

- 3.0** When performing computations on sparse matrices, latency in the memory hierarchy becomes much more of a factor. Sparse matrices lack the spatial locality in the data stream typically found in matrix operations. As a result, new matrix representations have been proposed.

One of the earliest sparse matrix representations is the Yale Sparse matrix Format. It stores an initial sparse $m \times n$ matrix, M in row form using three one-dimensional arrays. Let R be the number of nonzero entries in M . We construct an array A of length R that contains all nonzero entries of M (in left-to-right top-to-bottom order). We also construct a second array IA of length $m + 1$ (i.e., one entry per row, plus one). $IA(i)$ contains the index in A of the first nonzero of element of row i . Row i of the original matrix extends from $A(IA(i))$ to $A(IA(i+1)-1)$. The third array, JA , contains the column index of each element of A , so it also is of length R .

```

Row 0  [0, 0, -1, 0, 0, 4, 0, 0, 0, 7]
Row 1  [0, 2, 0, 3, 0, 0, 0, 0, 10, 0]
Row 2  [0, 0, 5, 8, 0, 0, 0, 0, 0, 9]
Row 3  [0, 0, 0, 0, 1, -2, 0, 0, 0, 0]
Row 4  [0, 6, 0, 0, 0, 0, 0, 0, 0, 0]
Row 5  [0, 0, 4, 0, 9, 0, 0, -2, 0, 0]
Row 6  [1, 2, 0, 0, 0, 0, 0, 0, -3, 0]
Row 7  [0, 11, 0, 0, 0, 0, 14, 0, 0, 0]
Row 8  [0, 6, 0, 11, -3, 0, 0, 0, 0, 0]
Row 9  [0, 0, 9, 0, 0, 0, 0, 7, 0, 0]
Row 10 [5, 0, 0, 0, 0, 0, 0, 0, -4, 0]
Row 11 [0, 1, 0, 0, 0, -1, 0, 0, 0, 0]

```

In terms of storage space, assuming that each element in matrix X is single precision floating point, compute the amount of storage used to store the Matrix above in Yale Sparse Matrix Format.